bigPint: A Bioconductor package that makes big data pint-sized

Lindsay Rutter[1][*] and Dianne Cook[2]

[*]Correspondence:
lindsayannerutter@gmail.com
[1] Bioinformatics and
Computational Biology Program,
Iowa State University, Ames, USA
Full list of author information is
available at the end of the article

Abstract

Keywords: sample; article; author

Background

Interactive data visualization is increasingly imperative in the biological sciences [1]. When performing RNA-seq studies, researchers wish to determine which genes are differentially expressed between treatment groups. Interactive visualization can help them assess differentially expressed gene (DEG) calls before performing any subsequent functional enrichment analyses. New visualization tools for genomic data have incorporated interactive capabilities, and some believe this trend could enhance the exploration of genomic data in the future [2]. Despite the growing appreciation of the inherent value of interactive graphics, the availability of effective and easy-to-use interactive visualization tools for RNA-seq data remains limited.

Some interactive visualization tools for genomic data have limited usability because they are only available on certain operating systems and/or require payment [3, 4, 5]. These limitations can be removed when tools are published on open-source repositories. Indeed, the Bioconductor project aims to foster interdisciplinary scientific research by promoting transparency and reproducibility while allowing software content to be used on Windows, MacOS, and Linux [6]. Bioconductor software is written in the R programming language, which also provides statistical and visualization methods that can facilitate the development of robust graphical tools in computational biology. Several interactive visualization methods for genomic data have been developed using Shiny, which is also based on R programming language [7, 8, 9].

We recently developed bigPint, an interactive data visualization software package available on Bioconductor. The bigPint package allows users to visually explore many types of large multivariate datasets, even though it was more specifically developed for RNA-seq data. In a recent methods paper, we used public RNA-seq datasets to demonstrate how bigPint graphics can help biologists detect crucial issues with normalization methods and DEG designation in ways not possible with numerical models [10]. We also applied bigPint visualization tools in a recent research paper that sought to elicit how nutrition and viral infection affect the honey bee transcriptome [11]. In the current paper, we explain the technical innovations and merits of the bigPint package, including new interactive visualization techniques that we believe can be helpful in the development and

usage of future biological visualization software. The bigPint website is located at https://bioconductor.org/packages/devel/bioc/html/bigPint.html and contains short vignette articles that introduce new users and provide suggested analysis pipelines, all written in reproducible code.

Results

Basic input

Each method in bigPint requires an input parameter data object. If a researcher is using the package to visualize RNA-seq data, then this data object should be a count table that contains the read counts for all genes of interest. The value in row i and column j should indicate how many reads have been assigned to gene i in sample j. This is the same input format required in popular RNA-seq count-based statistical methods, such as DESeq2, edgeR, limma, DSS, EBSeq, and BaySeq [?].

Several methods in bigPint also require an input parameter dataMetrics object. If a researcher is using the package to visualize RNA-seq data, then this dataMetrics object should be subset of the data (usually DEGs) where each case includes quantitative values of interest (such as fold change and FDR). This information can be easily derived from popular RNA-seq numerical analysis packages. Again, this conveniently allows users to work smoothly between visualizations in the bigPint package and models in other Bioconductor packages, abiding to the notion that the most efficient way to analyze large datasets is to iterate between models and visualizations.

Original features

1. Independent layers of interactivity

The Bioconductor community advanced the boundaries of biological visualization in the past and believes that interactive technology must be incorporated to continue these advancements in the future [6]. We will use the term *geom-drawing interactivity* to indicate user queries that draw geoms (graphical representation of the data, such as lines, hexagons, and points). This could mean the user adjusts sliders or selects buttons to draw a subset of data from the database as geoms (such as points). We will use the term *geom-manipulating interactivity* to indicate user interaction with already-drawn geoms. This could mean the user hovers over a geom (such as a hexagon) and obtains its associated metadata (such as the names of its contained genes). It could also mean the user zooming and panning to further alter how already-drawn geoms are displayed.

Our package introduces what we believe is a fairly new interactive visualization technology that is useful in the exploration of large biological datasets. Our technique allows for two independent layers of interactivity, for the foreground and background of the plot respectively. Each layer can include both geom-drawing and geom-manipulating interactivity. Our new technology can greatly enhance the exploration of large datasets, especially in cases where one layer contains large amounts of data (such as the full dataset) and the other layer contains smaller amounts of data (such as a data subset). Because the layers are independent, users can save time and computation by keeping the layer with more data drawn as is while only redrawing the layer with less data. We will now provide concrete examples of how this new two-layered interactivity method can improve upon several of the RNA-seq visualization tools in our package.

1a. Scatterplot matrices

Scatterplot matrices have appeared in statistical graphics literature for almost four decades and used across various fields of research [12, 13, 14, 15]. Previous user studies have shown that participants performed better when using animated rather than static versions of scatterplot matrices. Users also preferred animated scatterplot matrices and found them easier to understand as they can alleviate overplotting issues [16]. Interaction has been shown to extend the scatterplot matrix into an effective tool when representing large datasets [17]. We believe our new interactive visualization technology may further improve upon this long-standing plotting technique known for its effectiveness in exploratory multivariate data analysis.

1b. Litre plots

Problems still remain when scatterplot matrices are applied to large datasets. Physical space requirements grow exponentially by dimension size: for n-dimensional data, n^2 scatterplots are typically drawn. Hence, when extended to large dimensions, it becomes difficult to mentally link many small plots within the matrix [18]. Several techniques have been proposed to ameliorate this problem. Three dimensional scatterplots are useful but can cause occlusion and depth perception issues [18]. Other techniques like grand tours [19], projection pursuits [20, 21], and scagnostics [22] have been proposed.

Even though these alternative techniques are useful, they do not draw distributions across all cases (genes) and variables (samples). When analyzing RNA-seq data, we mainly want to compare replicate and treatment variability, which can be visually accomplished by plotting all genes and samples. It is useful to then superimpose DEGs to determine how they compare to the dataset as a whole. In light of this, we developed a plot that collapses the scatterplot matrix into one Cartesian coordinate system, allowing users to visualize all samples from one DEG of interest onto all samples of all genes in the dataset. We call this new plot a repLIcate TREatment ("litre") plot. An in depth explanation about the litre plot can be found in our previous methods paper [10].

We believe our interactive visualization technology is necessary for the litre plot. Drawing the background (all genes in the dataset) is the time-limiting step, whereas drawing the foreground (one DEG of interest) is immediate. Most users would like to quickly superimpose DEGs from a list one by one onto the background. This process would be unnecessarily time-prohibiting if the background needed to be redrawn each time the user superimposed the next DEG. Fortunately, our technology allows the user to immediately redraw the interactive foreground while the background remains unchanged but preserved in its interactive capabilities (hovering, zooming, and panning).

1c. Volcano plots

Volcano plots draw significance and fold change on the vertical and horizontal axes respectively. In RNA-seq studies, volcano plots allow users to check that genes were not falsely deemed significant due to outliers, low expression levels, and batch effects [23]. Researchers can benefit from the ability to quickly identify individual gene names in the volcano plot. This has been achieved with the identify() method

in R, which identifies the closest point in a scatterplot to the position clicked by the mouse [23]. We believe our interactive technology adds additional benefits to the volcano plot whereby users can filter on both statistics to immediately update the superimposed gene subset without having to redraw the more computationally-heavy background that contains all genes.

Useful features

Tailoring and saving static plots

Static plots can be saved as list objects in the R workspace and/or as JPG files to a directory chosen by the user. Saving plots into the R workspace allows users to integrate them into analysis workflows. It also allows them to tailor the plots (such as adding titles and changing label sizes) using the grammar of graphics via the conventional + syntax. Saving plots to a directory allows users to keep professional-looking files that can be inserted into proposals and talks. By default, the bigPint package saves static plots both in the R workspace and a directory (the default location is tempdir()).

Second feature layer

All plots, but static and interactive, allow for a selection of data to be plotted in a different way than the rest of the data. When analyzing RNA-seq data, this second feature layer could represent DEGs. There are three options for superimposing data subsets. First, users can declare data subsets using the previously-mentioned dataMetrics object by thresholding one of its quantitative variables. Second, users can simply declare a list of IDs to be overlaid using an object called geneList. This allows for more flexiblility than the first option. Third, the user can simply not overlay data subsets by leaving the dataMetrics and geneList objects with a value of NULL.

Group comparison filters

When users create static plots, the package automatically creates a separate plot for each pairwise combination of treatment groups from the inputted data. When users explore interactive plots, fields are dynamically generated from the inputted data so that any pairwise combination of treatment groups can be selected. Users can then quickly flip between different contrasts in their data. The package comes with an example dataset of soybean cotyledon data that has three treatment groups and examples of these contrast functionalities are shown on this data throughout the easy-to-follow articles on the package website.

Hexagonal binning

Users can create scatterplot matrices and litre plots using points (where each point represents one gene) or hexagonal binnings (where each hexagon color represents the number of genes in that area). Plotting each gene as a point allows for ideal levels of detail until overplotting occurs, which makes it difficult to determine how many genes are in given areas. In large datasets, this problem often persists even with techniques like alpha blending. Hexagonal binning has been used in prior software to successfully manage overplotting issues [17, 24] and has shown superior time

performance because less geometic objects need to be plotted. The bigPint package allows users to plot using either raw points or hexagonal binning, either of which can be useful depending on the dataset.

Hierarchical clustering

Users can conduct hierarchical clustering analyses on their data using the function plotClusters(). By default, the resulting clusters will be plotted as parallel coordinate lines superimposed onto side-by-side boxplots that represent the distribution of the full dataset. There are three main approaches in the plotClusters() function:

- Approach 1: The clusters are formed by clustering only on a user-defined subset of data (such as significant genes). Only these user-defined genes are overlaid as parallel coordinate lines.
- Approach 2: The clusters are formed by clustering the full dataset. Then, only a user-defined subset of data (such as significant genes) are overlaid as parallel coordinate lines.
- Approach 3: The clusters are formed by clustering the full dataset. All genes are overlaid as parallel coordinate lines.

The clustering algorithm is based on the hclust() and cutree() functions in the R stats package. It offers the same set of agglomeration methods ("ward.D", "ward.D2", "single", "complete", "average", "mcquitty", "median", and "centroid") with "ward.D" as the default. In many cases, users may want to save clusters derived from the plotClusters() function for later use, such as to overlay them onto scatterplot matrices, litre plots, and volcano plots. The gene IDs of each cluster can be saved as .RDS files for this purpose by setting the verbose option of the plotClusters() function to a value of TRUE.

Various plot aesthetics

Users can modify various aesthetics for both static and interactive plots, including hexagon size and point size. Some plots also allow for alpha blending, which has proven beneficial when plotting certain large datasets into parallel coordinate lines [25]. Statistical coloring is inconsistent in many numerous packages even though it can greatly enhance biological data visualization [26]. The bigPint package allows users to easily maintain consistent coloring across hierarchical clusters and when working between various plots.

Selection and aggregation

Some techniques that are effective in data exploration may lose their efficiency and eventually fail as data items increase in size. Two main approaches to solving these problems include data selection and data aggregation [27]. Data selection means that only a subset of the full dataset is displayed at a given time. The data subset can be selected through queries where data items are only displayed if they meet certain requirements [27]. The data subset can also be selected through interactive controls which allow the user to quickly examine between various subsets of the data [27]. Data aggregation means that the full dataset is divided into data subsets (called aggregates) that reduce the amount of data being visualized at once. When working with large datasets, an appropriate compromise is for users to be able to

perform both data selection and data aggregation [27]. The bigPint package allows users to easily perform data selection using queries and interactive controls (such as zooming, box and lasso selection, and panning) and to perform data aggregation using hierarchical clustering.

Shiny interactivity

Interactive plots in the bigPint package open as Shiny applications that consist of simple dashboards with "About" tabs that explain how to use the applications. They also include "Application" tabs that provide several input fields for the user to tailor their plots. Some of these input fields are generated dynamically from the inputted dataset so users have more flexibility in how they select any data subsets they wish to superimpose. In these applications, users can also download lists of selected genes and static images of interactive graphics to their local computers.

Shiny allows for linking between plots. Linking functionality between these plots plays a crucial role in rendering them suitable for large datasets [28, 29, 30]. By combining Shiny with Plotly and htmlwidgets functions, the bigPint package offers novel ways of dynamically and interactively working within and between plots.

Shiny applications can be launched on a local personal computer, hosted on a local or cloud-based server, or hosted for free on the shinyapps.io website. As such, interactive bigPint packages can be deployed on a personal computer using only a local file containing the data, the bigPint package and its dependencies, R / RStudio, and a browser recommended by Shiny (Google Chrome or Mozilla Firefox). This method does not require internet connectivity, which can be useful for users who are protecting sensitive data, analyzing or presenting data in contexts without reliable connectivity, or testing and developing applications.

Discussion

Researchers benefit when they are able to view multiple perspective of their data, especially when working with large datasets [31, 32]. The ability to select and aggregate data, threshold data to create subsets, link between multiple plots, interact with plots, and tailor various aesthetics in intelligent ways are all useful features of the bigPint package [1, 2, 18]. We expect that bigPint will enable researchers to generate and interact with intuitive, high quality, and reproducible plots from increasingly large datasets, including RNA-seq datasets.

Conclusion

Despite the growing appreciation of the inherent value in interactive graphics, the availability of easy-to-use and effective interactive exploratory visualization tools for RNA-seq data remains limited. In this paper, we introduced new visualization techniques with dynamic and interactive capabilities that are separately layered in the foreground and background of plots. We believe this methodology represents a fairly novel contribution to the field of interactive data visualization. Advocating state-of-the-art visualization tools is crucial for biology researchers to analyze and present their data and for visualization researchers to develop novel methods. Lessons learned from our open-source work may encourage the development of additional interactive visualization tools for various computational tasks.

Methods

bigPint was released under the GPL-3 license. Most bigPint visualization methods were constructed using htmlwidgets [33], ggplot2 [34], shiny [35], shinyapps.io [36], and plotly [37]. bigPint methods were tested on numerous RNA-seq datasets [10, 11]. The package website was constructed using the pkgdown software [38]. bigPint can be downloaded from the Bioconductor website [6].

Author details
[1] Bioinformatics and Computational Biology Program, Iowa State University, Ames, USA. [2] Econometrics and Business Statistics, Monash University, Clayton VIC, Australia.

References
1. O'Donoghue, S.I., Gavin, A.-C., Gehlenborg, N., Goodsell, D.S., Hériché, J.-K., Nielsen, C.B., North, C., Olson, A.J., Procter, J.B., Shattuck, D.W., *et al.*: Visualizing biological data—now and in the future. Nature methods **7**(3), 2 (2010)
2. Pavlopoulos, G.A., Malliarakis, D., Papanikolaou, N., Theodosiou, T., Enright, A.J., Iliopoulos, I.: Visualizing genome and systems biology: technologies, tools, implementation techniques and trends, past, present and future. Gigascience **4**(1), 38 (2015)
3. Seo, J., Shneiderman, B.: Interactively exploring hierarchical clustering results [gene identification]. Computer **35**(7), 80–86 (2002)
4. Ahlberg, C.: Spotfire: an information exploration environment. ACM SIGMOD Record **25**(4), 25–29 (1996)
5. Chu, L., Scharf, E., Kondo, T.: Genespringtm: tools for analyzing microarray expression data. Genome Informatics **12**, 227–229 (2001)
6. Gentleman, R.C., Carey, V.J., Bates, D.M., Bolstad, B., Dettling, M., Dudoit, S., Ellis, B., Gautier, L., Ge, Y., Gentry, J., *et al.*: Bioconductor: open software development for computational biology and bioinformatics. Genome biology **5**(10), 80 (2004)
7. Rue-Albrecht, K., Marini, F., Soneson, C., Lun, A.T.: isee: Interactive summarizedexperiment explorer. F1000Research **7** (2018)
8. Schultheis, H., Kuenne, C., Preussner, J., Wiegandt, R., Fust, A., Bentsen, M., Looso, M.: Wilson: Web-based interactive omics visualization. Bioinformatics **35**(6), 1055–1057 (2018)
9. Hughes, L.D., Lewis, S.A., Hughes, M.E.: Expressiondb: An open source platform for distributing genome-scale datasets. PloS one **12**(11), 0187457 (2017)
10. Rutter, L., Moran Lauter, A.N., Graham, M.A., Cook, D.: Visualization methods for rna-sequencing data analysis. Submitted to BMC Bioinformatics
11. Rutter, L., Carrillo-Tripp, J., Bonning, B.C., Cook, D., Toth, A.L., Dolezal, A.G.: Transcriptomic responses to diet quality and viral infection in apis mellifera. Submitted to BMC Genomics
12. Becker, R.A., Cleveland, W.S.: Brushing a scatterplot matrix: High-interaction graphical methods for analyzing multidimensional data. submitted for publication (1984)
13. Carr, D., Nicholson, W.: Graphical interaction tools for multiple 2-and 3-dimensional scatterplots. Technical report, Pacific Northwest Lab., Richland, WA (USA) (1984)
14. Tufte, E.R.: The Visual Display of Quantitative Information vol. 2. Graphics press Cheshire, CT, ??? (2001)
15. Tukey, P., Tukey, J.: Graphical display of data sets in three or more dimensions. Three papers in Interpreting Multivariate Data (ed. V. Barnett), 189–275. Chichester: Wiley (1981)
16. Chen, H., Engle, S., Joshi, A., Ragan, E.D., Yuksel, B.F., Harrison, L.: Using animation to alleviate overdraw in multiclass scatterplot matrices. In: Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems, p. 417 (2018). ACM
17. Carr, D.B., Littlefield, R.J., Nicholson, W., Littlefield, J.: Scatterplot matrix techniques for large n. Journal of the American Statistical Association **82**(398), 424–436 (1987)
18. Kerren, A., Ebert, A., Meyer, J.: Human-Centered Visualization Environments: GI-Dagstuhl Research Seminar, Dagstuhl Castle, Germany, March 5-8, 2006, Revised Papers vol. 4417. Springer, ??? (2007)
19. Asimov, D.: The grand tour: a tool for viewing multidimensional data. SIAM journal on scientific and statistical computing **6**(1), 128–143 (1985)
20. Friedman, J.H., Tukey, J.W.: A projection pursuit algorithm for exploratory data analysis. IEEE Transactions on computers **100**(9), 881–890 (1974)
21. Cook, D., Buja, A., Cabrera, J., Hurley, C.: Grand tour and projection pursuit. Journal of Computational and Graphical Statistics **4**(3), 155–172 (1995)
22. Wilkinson, L., Anand, A., Grossman, R.: Graph-theoretic scagnostics. In: IEEE Symposium on Information Visualization, 2005. INFOVIS 2005., pp. 157–164 (2005). IEEE
23. Li, W.: Application of volcano plots in analyses of mrna differential expressions with microarrays. arXiv preprint arXiv:1103.3434 (2011)

24. Harshbarger, J., Kratz, A., Carninci, P.: Deiva: a web application for interactive visual analysis of differential gene expression profiles. BMC genomics **18**(1), 47 (2017)
25. Unwin, A., Chen, C.-h., Härdle, W.: Computational Statistics and Data Visualization. Humboldt-Universität zu Berlin, Wirtschaftswissenschaftliche Fakultät, ??? (2007)
26. Yin, T., Cook, D., Lawrence, M.: ggbio: an r package for extending the grammar of graphics for genomic data. Genome biology **13**(8), 77 (2012)
27. Andrienko, G., Andrienko, N.: Blending aggregation and selection: Adapting parallel coordinates for the visualization of large datasets. The Cartographic Journal **42**(1), 49–60 (2005)
28. Nguyen, Q.V., Simoff, S., Qian, Y., Huang, M.L.: Deep exploration of multidimensional data with linkable scatterplots. In: Proceedings of the 9th International Symposium on Visual Information Communication and Interaction, pp. 43–50 (2016). ACM
29. Albuquerque, G., Eisemann, M., Lehmann, D.J., Theisel, H., Magnor, M.A.: Quality-based visualization matrices. In: VMV, pp. 341–350 (2009)
30. Heinrich, J., Stasko, J., Weiskopf, D.: The parallel coordinates matrix. EuroVis–Short Papers, 37–41 (2012)
31. Swayne, D.F., Lang, D.T., Buja, A., Cook, D.: Ggobi: evolving from xgobi into an extensible framework for interactive data visualization. Computational Statistics & Data Analysis **43**(4), 423–444 (2003)
32. Cook, D., Swayne, D.F., Buja, A.: Interactive and Dynamic Graphics for Data Analysis: with R and GGobi. Springer, ??? (2007)
33. Vaidyanathan, R., Xie, Y., Allaire, J., Cheng, J., Russell, K.: Htmlwidgets: HTML Widgets for R, 2016. https://cran.r-project.org/package=htmlwidgets Accessed 2018-12-20
34. Wickham, H.: Ggplot2: Elegant Graphics for Data Analysis. Springer, ??? (2016)
35. Chang, W., Cheng, J., Allaire, J., Xie, Y., McPherson, J.: Shiny: Web Application Framework for R [Computer Software]. https://cran.r-project.org/package=shiny Accessed 2018-12-20
36. RStudio: Integrated Development for R. http://www.rstudio.com Accessed 2018-12-20
37. Sievert, C., Parmer, C., Hocking, T., Chamberlain, S., Ram, K., Corvellec, M., Despouy, P.: Plotly: Create Interactive Web Graphics Via 'plotly.js'. https://cran.r-project.org/package=plotly Accessed 2018-12-20
38. Wickham, H., Hesselberth, J.: Pkgdown: Make Static HTML Documentation for a Package. https://cran.r-project.org/package=pkgdown Accessed 2018-12-20

Figures

Figure 1 Sample figure title. A short description of the figure content should go here.

Figure 2 Sample figure title. Figure legend text.

Tables

Table 1 Sample table title. This is where the description of the table should go.

	B1	B2	B3
A1	0.1	0.2	0.3
A2	...	..	.
A3	..	.	.

Additional Files
Additional file 1 — Sample additional file title
Additional file descriptions text (including details of how to view the file, if it is in a non-standard format or the file extension). This might refer to a multi-page table or a figure.

Additional file 2 — Sample additional file title
Additional file descriptions text.